



A Thesis on-

“AEGIS: A Constraint-Based Architecture for Safe LLM-Assisted Natural Language Analytics”

This thesis is submitted to the department of Computer Science & Engineering in partial fulfillment of the requirements for the degree of
Bachelor of Science in Computer Science & Engineering

Under the Course of-

Course Title: Thesis/Project Work
Course Code: CSE 4000(B)

Thesis Supervisor-

Name: Mst. Sahela Rahman
Designation: Lecturer

Prepared by-

Name: Md. Riaz
ID/Registration No: 0322310105101024
Batch: 16th
Semester: 8th
Session: Spring - 2023

**DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING
PUNDRA UNIVERSITY OF SCIENCE & TECHNOLOGY**

Date of submission: _____



A Thesis on-

“AEGIS: A Constraint-Based Architecture for Safe LLM-Assisted Natural Language Analytics”

This thesis is submitted to the department of Computer Science & Engineering in partial fulfillment of the requirements for the degree of
Bachelor of Science in Computer Science & Engineering

Signature of Supervisor

Signature of Student

**DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING
PUNDRA UNIVERSITY OF SCIENCE & TECHNOLOGY**

Date of submission: _____

CERTIFICATION OF ORIGINALITY

I hereby declare that this thesis titled "AEGIS: A Constraint-Based Architecture for Safe LLM-Assisted Natural Language Analytics" is my own original work, carried out under the supervision of Mst. Sahela Rahman, Lecturer, Department of Computer Science & Engineering, Pundra University of Science & Technology. To the best of my knowledge, this thesis contains no material previously published or written by another person, except where due reference is made in the text. This work has not been submitted, in whole or in part, for any other degree or qualification at this or any other institution.

Signature of Student

Md. Riaz

ID: 0322310105101024

CERTIFICATION OF APPROVAL

This is to certify that the research paper titled "AEGIS: A Constraint-Based Architecture for Safe LLM-Assisted Natural Language Analytics"

That it has been read, evaluated and accepted for fulfilment of the Academic Degree of Bachelor in Computer Science and Engineering (B.S.C. in CSE) at Pundra University of Science & Technology.

The acceptance decision is made based upon an independent review process that provides critically constructive feedback within a short turnaround time. This paper represents the author's independent work, critical analysis, and capability in undertaking literature research as per the academic policies and ethical standards of the university.

I certify that the candidate has completed the required research activities for the program and recommend that this piece of work as of acceptable standard for submission and for inclusion in the academic record of the University.

Mst. Sahela Rahman

Lecturer

Department of Computer Science and Engineering
Pundra University of Science & Technology

Date: _____

ACKNOWLEDGEMENT

I would like to express my sincere gratitude to my supervisor, Mst. Sahela Rahman, Lecturer, Department of Computer Science & Engineering, Pundra University of Science & Technology, for her continuous guidance, patience, and constructive feedback throughout the design, implementation, and evaluation of this research. Her insistence on precise, defensible claims shaped this thesis at every stage, from the formal safety proof to the honesty of its discussion of limitations.

I am also grateful to the faculty members of the Department of Computer Science & Engineering for the coursework and feedback that prepared me to undertake this research, and to my family for their patience and encouragement during the long implementation and evaluation cycles this thesis required.

Finally, I acknowledge the authors whose published research on natural language interfaces, text-to-SQL translation, and dashboard generation provided the foundation against which AEGIS's contributions are measured.

ABSTRACT

Relational databases hold data organizations need for decisions, but non-technical users often depend on developers to translate business questions into SQL. Natural-language interfaces try to close this gap, but many current systems still allow a language model to author executable SQL directly, making safety dependent on model behavior rather than system structure. This thesis presents AEGIS (Analytics Engine with Guaranteed Injection Safety), an architecture that restricts the model to intent extraction while deterministic software handles semantic mapping, permission enforcement, SQL compilation, validation, visualization selection, and widget persistence. AEGIS uses a closed semantic layer of approved metrics, dimensions, analytical patterns, and join paths, so the model never outputs SQL text. Instead, it produces a typed intent object that can be checked before any query is compiled. The final evaluation uses static nopCommerce datasets: a 500-question natural-language benchmark, 16 source-derived Admin analytics oracles, 80 Admin-fidelity phrasings, and focused semantic-coverage checks. On the 500-question live benchmark, AEGIS parsed 498 of 500 prompts, answered and executed 422 of 425 supported requests, and rejected or clarified 74 of 75 realistic boundary requests. Against the Admin analytics oracles, it achieved 16 of 16 execution validity, 16 of 16 shape accuracy, and 15 of 16 result accuracy. These results support the thesis's central claim that AEGIS is a bounded architecture for safe natural-language analytics over a governed semantic layer, not an unlimited text-to-SQL engine.

Table of Contents

Certification of Originality	i
Certification of Approval.....	ii
Acknowledgement	iii
Abstract	iv
Chapter 1: Introduction.....	1
1.1 Background.....	1
1.2 Problem Statement.....	1
1.3 Research Novelty and Motivation	2
1.4 Objectives and Contributions	3
Chapter 2: Literature Review and Research Gap	4
2.1 Natural Language Interfaces to Databases	4
2.2 Neural and LLM-Based Text-to-SQL	4
2.3 Natural Language for Visualization and Dashboards.....	5
2.4 Applied Conversational Business Intelligence	6
2.5 Comparative Summary	7
2.6 Research Gap Analysis.....	7
Chapter 3: Methodology	8
3.1 Research Paradigm	8
3.2 Formative Study of Reporting Patterns	9
3.3 Design Principles.....	10
3.4 Formal Model	11
3.5 Threat Model	11
3.6 System Architecture	12
3.7 Semantic Layer Design.....	14
3.8 Intent Parsing with Dynamic Vocabulary Injection	15
3.9 Safe Query Compiler	18

3.10 Visualization Selector	19
3.11 Widget Persistence and Reuse	20
Chapter 4: Experimental Work.....	21
4.1 Implementation	21
4.2 Experimental Environment.....	22
4.3 Benchmark Dataset Construction	22
4.4 Baseline and Oracle Comparisons	23
4.5 Evaluation Procedure.....	23
Chapter 5: Results and Discussion	25
5.1 Evaluation Overview	25
5.2 500-Question Natural-Language Benchmark.....	25
5.3 Admin Fidelity Benchmark	26
5.4 Focused Semantic Coverage.....	26
5.5 Safety Evaluation.....	27
5.6 Comparison With Direct LLM-to-SQL	27
5.7 Interpretation	28
Chapter 6: Limitations and Future Work.....	29
6.1 Limitations.....	29
6.2 Future Work.....	30
Chapter 7: Conclusion	30
References.....	32

List of Figures

Figure 1: Design Science Research workflow	8
Figure 2: AEGIS architecture pipeline	13
Figure 3: Semantic layer modularity	14
Figure 4: Vocabulary injection and coverage workflow	17
Figure 5: Two-layer SQL safety defence	19
Figure 6: Widget lifecycle and refresh model	20

List of Tables

Table 2.1: Comparative summary of related systems	7
Table 3.1: AEGIS reporting pattern taxonomy	9
Table 3.2: Semantic layer implementation contract	15
Table 3.3: AEGIS analytical patterns	18
Table 3.4: Visualization selector mapping	20
Table 4.1: Prototype module-to-architecture mapping	22
Table 4.2: Experimental setup	22
Table 4.3: Static evaluation datasets	23
Table 5.1: Current evaluation artifacts	25
Table 5.2: 500-question live benchmark results	26
Table 5.3: Admin analytics oracle benchmark	26
Table 5.4: Focused semantic coverage results	27
Table 5.5: Safety interpretation	27
Table 5.6: Structural comparison	28

CHAPTER 1

INTRODUCTION

1.1 Background

Relational databases store critical institutional data in organizations: financial records, customer accounts, sales transactions, and more. But accessing this data is uneven. Technical staff can write SQL queries to get any answer they need, while non-technical users have to wait for someone else to build them a report. This waiting is expensive. Analysis of enterprise reporting workflows shows that business users frequently wait days for new reports, and a recurring theme in institutional reporting is that many questions are variations of things already asked before: the same report with a different date range, or the same chart for a different department. These are not one-off questions; they are recurring reporting needs that should be served by saved, refreshable widgets rather than regenerated from scratch every time.

This thesis presents AEGIS (Analytics Engine with Guaranteed Injection Safety), a system that lets users describe their reporting needs in plain English and produces dynamic dashboard widgets that can be saved, refreshed, and reused as part of their daily workflow, without anyone writing SQL.

Natural language interfaces to databases (NLIDBs) try to solve this problem. The idea is simple: a user should be able to ask "which categories have the highest refund rates this month?" and get a correct, visual answer without writing SQL. Researchers have made good progress here through benchmarks such as Spider [1] and BIRD [2]. But there is still a gap between benchmark SQL generation and production-ready reporting, where answers must respect business definitions, permissions, safety rules, and reusable dashboard presentation.

1.2 Problem Statement

Three problems make up the gap between benchmark text-to-SQL accuracy and safe, production-ready natural-language reporting.

- **Safety:** Models may generate SQL directly, which can expose private data or create wrong joins.

- **Vocabulary mismatch:** Users ask with business terms such as "refund rate", while systems often expect column names.
- **No reusable widgets:** Most systems answer once and discard the result instead of saving refreshable dashboard widgets.

These problems are not primarily about building a smarter model; they are about designing the system properly around the model. AEGIS addresses this by splitting the work into stages. The LLM's only job is to understand what the user is asking and output a structured description of the request. Everything after that, matching to the right business terms, building the SQL, selecting the chart, and saving the widget, is done by fixed rules and pre-approved templates.

1.3 Research Novelty and Motivation

Existing text-to-SQL research asks: how accurately can a model generate SQL from natural language? This thesis asks a different question: how can large language models be used for language understanding while being structurally prevented from generating executable SQL at all? The two pipelines differ structurally.

Classical NL-to-SQL: natural-language request, then model-authored SQL, then query result.

AEGIS: natural-language request, then intent extraction, semantic constraint, deterministic compilation, and a safe analytical artifact.

The contribution of this thesis is therefore not improved SQL generation accuracy; it is constrained analytical artifact generation, a design approach that removes SQL generation from the LLM's role entirely and provides safety and semantic fidelity guarantees that no generative model can match unconditionally.

AEGIS does not propose a new LLM. The model is interchangeable across APIs that expose an OpenAI-compatible /v1/chat/completions interface. Because the LLM's only contract with the rest of the system is to produce a typed JSON intent object, model upgrades can improve language understanding without changing the compiler or safety infrastructure. This is model independence by design, not an implementation accident. The evaluation in Chapter 5 used an LLM API exposed through an OpenAI-compatible interface; the SQL compiler and safety layer did not depend on that provider.

1.4 Objectives and Contributions

The objective of this thesis is to design and evaluate AEGIS, a constraint-based architecture that lets an LLM understand reporting requests without allowing it to generate executable SQL. Query construction, validation, visualization, and widget persistence are handled by deterministic system components.

The specific objectives and contributions are:

1. A natural-language analytics architecture where the LLM is structurally prevented from writing executable SQL.
2. A closed semantic layer that maps business questions to approved metrics, dimensions, filters, analytical patterns, and join paths.
3. A deterministic query compiler that builds SQL from approved templates instead of model-written SQL.
4. A two-layer SQL safety mechanism combining structural prevention with post-compilation validation.
5. A widget-oriented workflow that turns natural-language answers into reusable dashboard artifacts.
6. A static nopCommerce evaluation corpus consisting of a 500-question natural-language benchmark, Admin analytics oracles, Admin-fidelity phrasings, and focused semantic-coverage checks.
7. An empirical evaluation showing broad supported-request coverage, strong boundary rejection, and a visible remaining Admin fidelity gap rather than a suspicious claim of perfect performance.

CHAPTER 2

LITERATURE REVIEW AND RESEARCH GAP

The literature review focuses on the work most directly comparable to AEGIS. It is organized as short review blocks so each source can be checked quickly by contribution, limitation, and the specific gap it leaves for this thesis.

2.1 Natural Language Interfaces to Databases

NLIDB surveys [3, 4]

- **Contribution:** Compare major natural-language database interface designs, including keyword, pattern, parsing, and grammar-based systems.
- **Limitations:** Safety is treated mostly as a secondary concern; even SQL-injection discussion ends with generic filtering advice.
- **Gap for AEGIS:** No reviewed system makes SQL safety a structural design property.

NaLIR [5]

- **Contribution:** Uses an intermediate query tree and asks the user to choose between possible interpretations.
- **Limitations:** Correctness depends on user disambiguation, and the final output remains executable SQL.
- **Gap for AEGIS:** Ambiguity is handled after parsing, but unsafe or unauthorized queries are not prevented by design.

Veezoo [6]

- **Contribution:** Uses an editable Knowledge Graph that maps business language to database concepts.
- **Limitations:** The semantic layer improves usability, but the paper does not evaluate SQL safety or permission enforcement.
- **Gap for AEGIS:** A semantic layer exists in prior work, but not as a safety boundary.

2.2 Neural and LLM-Based Text-to-SQL

Spider and BIRD [1, 2]

- **Contribution:** Provide large text-to-SQL benchmarks for complex and cross-domain database questions.
- **Limitations:** They measure SQL correctness and execution accuracy, not whether a query should be allowed.
- **Gap for AEGIS:** Benchmark success does not equal database safety or business-policy compliance.

RAT-SQL [7]

- **Contribution:** Models schema relations directly so table and column linking can improve on complex schemas.
- **Limitations:** Errors still frequently come from wrong table or column selection.
- **Gap for AEGIS:** Better schema encoding reduces some mistakes but does not create an authorization boundary.

PICARD [8]

- **Contribution:** Constrains decoding so generated SQL is more likely to be syntactically valid.
- **Limitations:** The model still authors the SQL, and grammar validity does not prove the query is semantically safe.
- **Gap for AEGIS:** A valid SQL string can still express the wrong intent or access the wrong data.

G-SQL and TriSQL [9, 10]

- **Contribution:** Show two strong directions: rule-guided schema-aware translation and multi-stage LLM refinement.
- **Limitations:** G-SQL has limited qualitative evaluation, while TriSQL still depends on LLM-authored executable SQL.
- **Gap for AEGIS:** The closest prior work improves accuracy, but does not remove SQL generation from the LLM output space.

2.3 Natural Language for Visualization and Dashboards

nl4dv and DataTone [11, 12]

- **Contribution:** Map natural-language questions into visualization tasks and expose ambiguity instead of silently guessing.
- **Limitations:** They focus on visualization specification, not protected database execution or persistent dashboard widgets.
- **Gap for AEGIS:** Visualization systems solve chart selection but leave SQL safety and reuse outside the architecture.

DashBot [13]

- **Contribution:** Uses deep reinforcement learning and dashboard-design rules to compose multi-chart dashboards.
- **Limitations:** It has no natural-language-to-SQL stage and does not handle permission-controlled production database access.
- **Gap for AEGIS:** Dashboard composition is useful, but it does not solve safe language-to-data translation.

2.4 Applied Conversational Business Intelligence

Conversational BI assistants [14]

- **Contribution:** Demonstrate that LLM-based chat can be connected to dashboards and SQL tools.
- **Limitations:** Direct SQL execution through an LLM tool loop creates a large attack surface and is rarely evaluated adversarially.
- **Gap for AEGIS:** A production assistant needs a constrained execution layer, not only a conversational interface.

Explanatory dashboard analytics [15]

- **Contribution:** Shows that deterministic computation followed by LLM narration can improve reliability in business explanation tasks.
- **Limitations:** The work targets dashboard explanation, not safe SQL compilation or reusable widget creation.
- **Gap for AEGIS:** It supports AEGIS's choice to let deterministic code compute results while the LLM handles language.

2.5 Comparative Summary

Table 2.1: Comparative summary of the most closely related systems.

System	NL	Semantic	Safe SQL	Visual	Widget	Coverage	Evaluation
Spider/BIRD	Yes	-	-	-	-	-	Benchmark
RAT-SQL/PICARD	Yes	-	Partial	-	-	-	Benchmark
G-SQL/TriSQL	Yes	Partial	Partial	-	-	-	Benchmark
NaLIR	Yes	-	-	-	-	-	User study
Veezoo	Yes	Yes	-	-	-	-	User study
nl4dv/DataTone	Yes	-	-	Yes	-	-	User study
DashBot	-	-	-	Yes	Partial	-	Synthetic
Conversational BI	Yes	-	-	Yes	-	-	Demo
AEGIS	Yes	Yes	Yes	Yes	Yes	Yes	Production

2.6 Research Gap Analysis

- **Safety gap:** Most systems measure answer accuracy, but do not evaluate unsafe SQL behavior.
- **Semantic-layer gap:** Semantic layers appear in prior work, but mainly for usability and matching, not as execution boundaries.
- **Visualization gap:** Visualization systems produce charts, but usually operate outside permission-controlled SQL execution.
- **Persistence gap:** Prior tools often answer one question at a time, while recurring business reports need refreshable outputs.
- **Thesis position:** AEGIS addresses these gaps through a bounded vocabulary, deterministic query compiler, safe visualization selector, and reusable widget model.

CHAPTER 3

METHODOLOGY

3.1 Research Paradigm

This thesis follows a Design Science Research paradigm because the main contribution is a built and evaluated artifact: the AEGIS system. The paradigm is summarized through three requirements:

- **Problem relevance:** Representative reporting requests motivate the need for the artifact.
- **Design structure:** The semantic layer, threat model, and compiler boundaries define how the artifact is designed.
- **Evaluation:** Static nopCommerce datasets evaluate supported-request coverage, boundary rejection, Admin fidelity, execution validity, and remaining implementation limits.

The artifact was refined through three build-evaluate cycles:

- **Cycle 1:** Built the initial semantic layer and compiler, then tested the core safety path.
- **Cycle 2:** Expanded coverage to the eleven analytical patterns and replaced manual synonym handling with vocabulary injection.
- **Cycle 3:** Added widget persistence and expanded the benchmark execution harness.




Figure 1: Design Science Research workflow for AEGIS

3.2 Formative Study of Reporting Patterns

The eleven-pattern taxonomy used throughout this thesis (Table 3.3) was derived from a design-time review of representative e-commerce and administrative reporting requests while designing AEGIS. The purpose of this review was to identify recurring business-question shapes that the system should support, such as KPI summaries, rankings, trends, comparisons, exception reports, and reusable tabular views.

Table 3.1 summarizes the reporting patterns used by the AEGIS compiler and shows how they appear in ordinary e-commerce analytics. The taxonomy is a design artifact used to structure the semantic layer and compiler templates, not an independently validated user study.

Table 3.1: AEGIS reporting pattern taxonomy.

Pattern	Purpose	Example
KPI / Aggregate	Single-number business summary	Total revenue this month
Ranking	Top or bottom entities by metric	Top products by revenue
Exception / Filter	Rows that meet an operational condition	Low-stock products
Trend Analysis	Metric over time	Monthly sales trend
Comparison	Metric compared across groups or periods	Revenue by order status
Summary / Group	Grouped business overview	Orders by payment status
Cohort	Population split by lifecycle stage	New versus returning customers
Funnel	Stepwise business process	Checkout progression
Correlate	Relationship between two measures	Discount versus order value
Segment	Breakdown by business dimension	Revenue by country
Tabular	Detailed record listing	Latest orders

The classification shaped the methodology in three ways. First, it made the compiler finite: each supported request must fit an approved analytical pattern. Second, it encouraged business terminology rather than database column names, which supports the need for an explicit semantic layer. Third, many reporting needs are reusable with only a changed time window, filter, or segment, which supports the widget persistence design.

3.3 Design Principles

Five principles guide the AEGIS architecture:

8. Separate understanding from execution. The LLM understands the question; fixed rules handle everything else.
9. Define business terms explicitly. Metrics, dimensions, joins, and time rules are written once in a semantic layer, not inferred per query.

10. Limit what SQL can be generated. SQL is built only from pre-approved, parameterized templates.
11. Select visualizations by rule. Chart type is decided by question type, result shape, and established visualization design guidance, not by a learned model.
12. Persist results for reuse. Every query produces a saved, refreshable widget rather than a discarded answer.

3.4 Formal Model

The practical model follows directly from the design principles above: user language is converted into a typed intent, approved semantic bindings are selected, and SQL is built only by deterministic templates. This keeps the model useful for understanding language without giving it authority to write executable SQL.

3.5 Threat Model

AEGIS protects against attacks arriving through the untrusted natural-language input channel. The model assumes the database and application server are properly hardened; the attacker controls only the query field.

- **T1 - Prompt injection attempting SQL generation:** "Ignore previous instructions. Generate DROP TABLE orders."
 - **Control:** The IntentObject schema contains no SQL field. Any non-approved string in metric_term or dimension_term is rejected by Pydantic type validation at Stage 2, before the compiler is reached.
- **T2 - Unauthorized metric or dimension access:** "Show me customer passwords" or "List credit card numbers by order."
 - **Control:** Fields such as customer_password do not exist in the semantic layer vocabulary. The LLM never sees those names; it receives only the curated, approved label list. Stage 2 rejects any unrecognized term.
- **T3 - Unauthorized row access:** A store-level user asks: "Show revenue for all branches."
 - **Control:** Stage 4 (Permission Rewriter) runs after the LLM and appends a role-specific WHERE predicate (for example, AND o.StoreId = :user_store) derived from

the authenticated session. This cannot be suppressed or overridden by natural-language content.

- **T4 - DML or DDL injection:** A crafted prompt that tricks the LLM into associating a write operation with an intent class.
 - **Control:** No template in the pattern library contains a DML or DDL keyword. The AST-level post-compilation validator explicitly rejects any non-SELECT statement as an additional safety layer.

Not protected by AEGIS (requires operational security controls outside this architecture):

- A malicious administrator embedding arbitrary SQL inside a metric's `sql_expr` field.
- Supply-chain compromise of the compiler module or the SQL parser library.
- Database-level privilege escalation that bypasses the application layer.
- LLM provider infrastructure compromise or model poisoning.

Explicitly documenting out-of-scope threats is itself a contribution: prior NL-to-SQL work rarely specifies the boundary of its safety claims, which makes meaningful security comparison difficult.

3.6 System Architecture

Every user query travels through exactly seven stages. Stages 2 through 7 contain no artificial intelligence; they are deterministic code. Rejection at any stage produces a structured clarification message rather than a best-effort guess.

- **Stage 1 - Intent Extraction:** A lightweight LLM reads the query together with a system prompt built from the semantic layer, and outputs a validated IntentObject (intent class, metric term, dimension term, filters, sort, limit, confidence). This is the only stage that involves artificial intelligence.
- **Stage 2 - Coverage Validation:** The server checks that both the metric term and the dimension term exist in the semantic layer vocabulary before anything else runs. Unknown identifiers are rejected here, with a structured message listing the available identifiers, rather than being passed to the compiler.
- **Stage 3 - Semantic Mapping:** Business-logic aliases are expanded (for example, 'abandoned' maps to a specific OrderStatusId), and relative time expressions such as 'this month' are resolved to concrete date predicates.
- **Stage 4 - Permission Rewriting:** A role-specific WHERE predicate is appended based on the authenticated user's session. This runs after the LLM has already finished, so no natural-language content can influence it.
- **Stage 5 - SQL Compilation:** A breadth-first search over the join graph finds the minimal join path connecting the tables required by the resolved metric and dimension, and pre-compiled SQL expressions are substituted into a parameterized template. No SQL text is ever assembled from concatenated user input.
- **Stage 6 - Visualization Selection:** A rule engine maps the intent class and result shape to a default chart type. This stage contains no learned model.
- **Stage 7 - Widget Persistence:** The analysis plan is hashed (SHA-256) to detect duplicates, and the query, chart configuration, and access rules are stored as a widget artifact that can be refreshed on a schedule.




Figure 2: AEGIS architecture pipeline (User Request to Dashboard Widget)

3.7 Semantic Layer Design

The semantic layer is the most important non-AI component of AEGIS. It separates business language from the underlying database structure and defines exactly which metrics, dimensions, joins, predicates, and permissions are allowed to exist. The semantic layer is also the main per-deployment implementation surface: to support another system, the developer defines that system's governed business vocabulary and join paths while preserving the same AEGIS architecture.

This presentation follows the style used by related systems: Veezoo describes a Knowledge Graph, parser, query processor, and visualization engine [6]; G-SQL describes JSON schema serialization and rule-guided clause construction [9]; and NL4DV exposes a JSON analytic specification for attributes, tasks, and visualization choices [11]. AEGIS therefore defines the semantic layer as an explicit implementation contract rather than as a general idea.




Figure 3: Semantic layer modularity - composable blocks vs. free-form SQL generation

Table 3.2: Semantic layer implementation contract.

Object	Required fields	nopCommerce example
Metric	id, label, description, sql_expr, binding_table, required_joins, time_anchor	revenue -> SUM(COALESCE(o.OrderTotal,0)), bound to Order
Grain rule	item_grain_equivalent when an order-level metric is grouped by item data	revenue by product uses line_item_revenue instead of duplicating order totals
Dimension	id, label, sql_expr, binding_table, datatype, entity, group_expr	category_name -> c.Name, reached through Product_Category_Mapping and Category
Predicate	label, SQL predicate, parameter field, datatype	payment_status -> o.PaymentStatusId = :payment_status
Time anchor	metric-specific date column for period filters	customer_count uses cu.CreatedOnUtc; order_count uses o.CreatedOnUtc
Join path	source table, target table, ON clause	Order -> OrderItem -> Product -> Category
Mandatory predicate	always-on table rule appended by the compiler	Order, Product, and Customer include Deleted = 0

A compact example of the implementation contract is shown below:


```

Metric(
  id="revenue",
  sql_expr="SUM(COALESCE(o.OrderTotal, 0))",
  binding_table="Order",
  time_anchor="o.CreatedOnUtc",
  item_grain_equivalent="line_item_revenue")

Dimension(
  id="category_name",
  sql_expr="c.Name",
  binding_table="Category",
  required_joins=["Product_Category_Mapping", "Category"])

```

In the nopCommerce prototype, the semantic layer defines the governed metrics, dimensions, predicates, and join paths needed for the evaluated e-commerce analytics scope. The full nopCommerce schema is larger than this exposed vocabulary. Tables and fields that are not represented in the semantic layer cannot be requested through AEGIS, which is how the architecture keeps the answerable space useful but finite.

3.8 Intent Parsing with Dynamic Vocabulary Injection

The central technique that makes Stage 1 reliable is vocabulary injection. At startup, AEGIS builds the system prompt by listing every approved metric and dimension name, with a plain-English description, directly from the semantic layer (approximately 1,100 tokens for 15 metrics and 34 dimensions). The LLM sees exactly which identifiers are valid and maps any user phrasing to the correct identifier without a manually maintained synonym list. This has three advantages over a synonym dictionary: zero maintenance, since adding a metric automatically updates the vocabulary the model sees; broad coverage of arbitrary user phrasing, since the model performs the mapping rather than a fixed lookup table; and token efficiency, since the full vocabulary for 15 metrics and 34 dimensions fits in roughly 1,100 tokens.

Structured output enforcement for LLMs [16] constrains the model's response to a fixed JSON schema before any downstream validation occurs, which is why Stage 1 can rely on a typed `IntentObject` rather than free-form text: malformed or off-schema output is rejected at the API boundary, before Stage 2 coverage validation ever runs.

The output schema enforces typed fields:

```

{
  "intent_class": "kpi | ranking | trend | comparison | exception |
    summary | segment | funnel | cohort | correlate | tabular",
  "metric_term": "string",

```

```

"dimension_term": "string or null",
"time_term": "string or null",
"filters": [{"field": "string", "operator": "string", "value": "string"}],
"sort": "asc | desc | null",
"limit": "integer or null",
"confidence": "low | medium | high",
"needs_clarification": "boolean"
}

```




Figure 4: Vocabulary injection and original-question coverage workflow

3.9 Safe Query Compiler

The compiler instantiates SQL from a library of parameterized templates, one per analytics pattern. Its role is not to be creative; its role is to prove that a grounded plan can be rendered using only semantic-layer objects.

The compiler procedure is deterministic:

Input: AnalysisPlan(pattern, metric, dimension, filters, time_range)

1. collect required tables from metric, dimension, and declared required_joins
2. replace order-grain metrics with declared item-grain equivalents when needed
3. resolve the minimal join path with BFS over the semantic join graph
4. assemble SELECT from approved sql_expr values only
5. assemble WHERE from normalized time_range and governed predicates
6. append mandatory table predicates such as Deleted = 0
7. add GROUP BY, ORDER BY, and LIMIT according to the analytical pattern
8. reject the final SQL if it contains forbidden constructs

Output: read-only SQL string, bound parameters, rationale log

The important implementation detail is that user text never enters a SQL identifier position. Identifiers come from Metric and Dimension objects; literal values are carried as parameters; and join clauses come from the join graph. If any slot cannot be grounded, the resolver returns reject or clarify rather than allowing the compiler to guess.

Table 3.3: The eleven AEGIS analytical patterns.

Pattern	Required slots	Optional slots	Default visual
KPI (Aggregate)	metric	time_rule, filter	kpi_card
Ranking	metric, dimension	time_rule, filter, limit	bar_chart
Trend	metric, time_grain	time_rule, filter	line_chart
Comparison	metric, segment	time_rule, filter	grouped_bar
Exception	metric, threshold	dimension, time_rule	table
Summary	metric[], dimension	time_rule, filter	multi_card
Segment	metric, dimension	time_rule, filter	pie_chart
Funnel	metric, stages	time_rule, filter	funnel_chart
Cohort	metric, group_def	time_rule, filter	grouped_bar
Correlate	metric, attribute	time_rule, filter	scatter_plot
Tabular	dimension	filters, time_rule	table

Two safety layers apply in sequence. Layer 1 is a parameterized query engine that separates SQL structure from user-supplied inputs, so no user text is ever concatenated into the SQL string. Layer 2 is a post-compilation safety scanner that rejects any assembled query containing a forbidden construct: non-SELECT statements, UNION, EXCEPT or INTERSECT, EXEC, or references to system tables. If any forbidden pattern is detected, the compiler raises a SecurityError rather than returning a partially safe query.




Figure 5: Two-layer SQL safety defence

3.10 Visualization Selector

Table 3.4: Visualization selector mapping.

Intent	Result shape	Selected visualization
KPI	scalar	KPI card
Ranking	1 measure, up to 20 categories	Horizontal bar chart
Trend	1 measure, time series	Line chart
Comparison	1 measure, 2-4 segments	Grouped bar chart
Exception	row-level detail	Sortable table
Summary	2-4 scalar measures	KPI card grid
Segment	1 measure, categorical	Pie chart
Funnel	ordered conversion stages	Funnel chart
Cohort	1 measure, 2+ groups	Grouped bar chart
Correlate	2 measures, continuous	Scatter plot
Tabular	raw records	Sortable table

Two additional rules apply after the data is known, rather than at selection time: bar charts with more than 20 categories are automatically converted to a table, and pie charts with more than 8 slices are converted to a bar chart. Both rules exist because the initial chart choice is made before the exact cardinality of the result is known.

3.11 Widget Persistence and Reuse

Each widget stores a unique identifier (a SHA-256 hash of the analysis plan), the original question, the analysis plan in JSON form, a hash of the SQL template used, chart configuration, timestamps, access rules, and run history. A new question triggers the full seven-stage pipeline; if an identical widget already exists, determined by a match on the plan hash, the cached artifact is returned immediately rather than recompiled. Scheduled refresh re-executes the stored SQL against fresh data on a configurable interval, which directly addresses the design-time observation that reporting requests are often recurring rather than one-off: a widget answers the question once and then continues answering it as new data arrives, rather than requiring the same natural-language request to be re-processed from scratch every time.




Figure 6: Widget lifecycle and refresh model

CHAPTER 4

EXPERIMENTAL WORK

4.1 Implementation

AEGIS is implemented as a web application with a vanilla HTML and JavaScript frontend and a Python FastAPI backend, targeting a nopCommerce 4.70-style e-commerce schema. The prototype follows the architecture from Chapter 3: the model extracts intent, the semantic layer supplies governed business definitions, and deterministic compiler templates produce SQL.

- **LLM integration:** An LLM API exposed through an OpenAI-compatible `/v1/chat/completions` interface with structured JSON output enforcement. The prompt injects approved metric and dimension identifiers at request time. The compiler and safety layer are provider-independent because they consume only the typed intent object, not model-written SQL.
- **Semantic layer:** Python configuration modules define governed metrics, dimensions, predicates, time anchors, join paths, grain rules, and mandatory platform filters.
- **SQL compiler:** Parameterized MySQL templates build read-only SQL from approved analytical patterns. The compiler logs the selected pattern, resolved tables, join path, metric expression, dimension expression, and safety result.
- **Visualization selector:** Rule-based chart selection maps analytical patterns and result shapes to dashboard widgets.
- **Widget engine:** The prototype stores widget metadata locally with plan-hash deduplication. The interface is designed so production storage can be moved to a database.
- **Coverage validator:** Coverage analysis checks the original user question against the semantic layer so unsupported concepts can be rejected or clarified instead of silently mapped to the nearest available metric.

Table 4.1: Prototype module-to-architecture mapping.

Architecture stage	Implementation module	Main technical responsibility
Intent extraction	intent_parser.py, models.py	OpenAI-compatible structured JSON output into IntentObject
Grounding	grounding.py, mapper.py	Ranked semantic-layer binding with resolved, ambiguous, unsupported, or absent outcomes
Coverage	coverage.py	Original-question inspection for unsupported e-commerce concepts
Semantic layer	semantic_layer.py	Metrics, dimensions, predicates, time anchors, join graph, and mandatory filters
Compilation	compiler.py	Pattern templates, BFS join-path resolution, parameter binding, and SQL safety scan
Visualization	visualization.py	Rule-based chart selection from pattern and result shape
Persistence	widget_engine.py	Plan hashing, widget storage, refresh metadata, and reuse

This module map is included to make the prototype auditable: the thesis architecture is not only a conceptual pipeline, and each stage has a corresponding implementation boundary.

4.2 Experimental Environment

Table 4.2: Experimental setup.

Setup item	Configuration
Database engine	MySQL 8.0
Execution mode	Local/Laragon or Docker-compatible MySQL evaluation database
Application schema	nopCommerce 4.70-style e-commerce schema
Dataset scope	Orders, customers, products, categories, manufacturers, payments, shipping, refunds, stores, countries, and search terms
LLM interface	OpenAI-compatible /v1/chat/completions API
Evaluation scope	Single-domain prototype evaluation over nopCommerce analytics

4.3 Benchmark Dataset Construction

This evaluation is a prototype evaluation, not a large-scale independent leaderboard study. Its goal is to test whether the AEGIS architecture provides safe, governed natural-language analytics for a realistic e-commerce schema. General text-to-SQL benchmarks such as Spider and BIRD are valuable for SQL generation research, but they do not measure reusable dashboard widgets, semantic layer governance, or refusal of unsupported business questions.

Table 4.3: Static evaluation datasets.

Dataset	Size	Purpose
Natural-language benchmark	500 questions	Tests broad nopCommerce semantic-layer coverage and realistic boundary refusal.
Admin analytics oracles	16 tasks	Checks fidelity against source-derived nopCommerce Admin reporting logic.
Admin-fidelity phrasings	80 prompts	Tests five natural phrasings for each Admin oracle task.
Semantic coverage suite	25 checks	Tests representative supported combinations and focused boundary cases.

All benchmark datasets are static repository artifacts. They are not regenerated during evaluation, because a thesis dataset must be inspectable and stable. The reported result files are also stored with the repository so a reader can compare the manuscript tables with the underlying evidence.

4.4 Baseline and Oracle Comparisons

The evaluation uses two kinds of comparison. First, the thesis discusses the structural contrast between AEGIS and direct LLM-to-SQL systems, where the model is allowed to generate SQL text. Second, the Admin fidelity benchmark compares AEGIS output with source-derived nopCommerce Admin analytics oracles. The second comparison is the stronger evidence for result accuracy because it checks returned business values rather than only whether SQL text was emitted.

- **Direct LLM-to-SQL:** The model is prompted to generate SQL directly from a database schema. This baseline has broad expressive freedom but weak governance because joins, filters, and business definitions are inferred per request.
- **Admin oracle comparison:** Expected outputs are extracted from nopCommerce Admin analytics logic. AEGIS can use different SQL text, but the returned result shape and values must match the oracle.

4.5 Evaluation Procedure

The current evaluation focuses on measurements that can be reproduced from static benchmark data, verifier scripts, and recorded result artifacts:

- **RQ1:** How reliably does the LLM parser produce typed reporting intents on the 500-question natural-language dataset?
- **RQ2:** Does AEGIS answer supported semantic-layer questions while rejecting or clarifying realistic unsupported e-commerce questions?
- **RQ3:** Does the compiled SQL execute successfully against the evaluation database?
- **RQ4:** Do AEGIS outputs match the shape and values of source-derived nopCommerce Admin analytics oracles?
- **RQ5:** What implementation gaps remain after the current semantic-layer and compiler updates?

CHAPTER 5

RESULTS AND DISCUSSION

This chapter reports the current AEGIS prototype evaluation using static nopCommerce datasets committed with the repository. The evaluation separates four questions that are often mixed together: whether the system can parse a natural-language request, whether the request is covered by the semantic layer, whether the compiled SQL executes, and whether the returned result matches the expected business answer.

The results should not be read as a claim that AEGIS is a perfect analytics system. A perfect score on every table would be suspicious for a prototype. Instead, the evaluation shows where the architecture is strong, where the current nopCommerce implementation is already useful, and where a visible implementation gap remains.

5.1 Evaluation Overview

Table 5.1: Current evaluation artifacts.

Artifact	Purpose	Status
500-question natural-language dataset	Tests broad user-facing e-commerce questions over the implemented semantic layer.	425 supported requests and 75 realistic boundary requests.
Admin analytics oracles	Compares AEGIS outputs with source-derived nopCommerce Admin reporting logic.	16 oracle tasks.
Admin-fidelity phrasings	Checks whether the same Admin task can be requested in multiple natural phrasings.	80 prompts, five per oracle task.
Focused semantic coverage suite	Checks representative metric, dimension, trend, ranking, exception, and boundary cases.	20 supported and 5 boundary checks.

These datasets are static research artifacts, not generated at evaluation time. This matters because a thesis result must be reproducible: another reader should be able to inspect the exact questions, oracle definitions, scripts, and result files used for the reported claims.

5.2 500-Question Natural-Language Benchmark

The main breadth benchmark contains natural-language questions that a store owner or administrator might ask about orders, revenue, customers, products, refunds, inventory, stores, countries, payment status, shipping, and search terms. The dataset intentionally includes unsupported questions as well, because a governed analytics system must know when not to answer.

Table 5.2: 500-question live benchmark results.

Metric	Result	Interpretation
Parser success	498/500 (99.6%)	The LLM API usually returned a valid typed intent.
Supported intent exact match	345/425 (81.2%)	Many prompts were parsed exactly; remaining supported prompts often still compiled to a valid answer.
Supported answer rate	422/425 (99.3%)	AEGIS answered almost all requests that were inside the semantic layer.
Supported execution validity	422/425 (99.3%)	The answered supported requests also executed successfully.
Boundary rejection accuracy	74/75 (98.7%)	Most realistic e-commerce questions outside the semantic layer were rejected or clarified.

This result supports the bounded-system claim. AEGIS is not trying to answer infinite arbitrary SQL questions. It is trying to answer a large, useful set of questions made available by the semantic layer and to avoid silently inventing answers for unavailable concepts.

5.3 Admin Fidelity Benchmark

The Admin fidelity benchmark is stricter than asking whether AEGIS can produce some chartable SQL. It extracts reporting expectations from nopCommerce Admin analytics and compares the returned shape and values against those source-derived oracles. Exact SQL text is not required; the important question is whether the answer has the correct business meaning.

Table 5.3: Admin analytics oracle benchmark.

Metric	Result	Interpretation
Execution validity	16/16 (100.0%)	Every Admin oracle task produced executable SQL.
Shape accuracy	16/16 (100.0%)	Every output used the expected row/column structure.
Result accuracy	15/16 (93.8%)	One dashboard order-average matrix still differs from nopCommerce.

The remaining mismatch is useful evidence rather than an embarrassment. It shows that the benchmark is capable of finding a real gap. The gap should be fixed by adding a general multi-period matrix-summary primitive to the compiler, not by hardcoding a nopCommerce report shortcut.

5.4 Focused Semantic Coverage

Table 5.4: Focused semantic coverage results.

Metric	Result	Interpretation
Supported execution validity	20/20 (100.0%)	Representative supported requests executed.
Supported shape accuracy	20/20 (100.0%)	Outputs matched the expected widget form.
Supported result accuracy	20/20 (100.0%)	Focused supported cases matched expected values.
Boundary rejection accuracy	5/5 (100.0%)	Focused unsupported cases were rejected.

This small suite is not presented as proof of perfection. It is a regression and coverage check over representative semantic-layer combinations. The broader 500-question benchmark and the Admin oracle mismatch are what keep the evaluation honest.

5.5 Safety Evaluation

The central architectural safety property is that the language model does not write SQL. It returns a typed intent object. SQL is then produced by deterministic templates over approved semantic-layer definitions and checked before execution. Therefore, prompt text from the user is not interpolated into executable SQL.

Table 5.5: Safety interpretation.

Risk	AEGIS control	Remaining boundary
User asks for a write operation	Intent schema and compiler templates do not include write primitives.	The system must still explain refusal clearly to the user.
User uses unsupported business terms	Coverage analysis checks the original user question against the semantic layer.	Coverage is finite and depends on deployment configuration.
LLM misunderstands the request	Compiled SQL remains structurally safe.	The answer may still be semantically wrong if intent extraction is wrong.

5.6 Comparison With Direct LLM-to-SQL

A direct LLM-to-SQL baseline can appear attractive because it can always try to write a SELECT statement. However, allowing the model to author SQL moves business definitions, join choices, security behavior, and dialect correctness into a probabilistic component. AEGIS deliberately gives up unlimited query flexibility in exchange for governed definitions, deterministic compilation, and an explicit refusal path.

Table 5.6: Structural comparison.

Property	Direct LLM-to-SQL	AEGIS
SQL generation	Model writes SQL text	Compiler expands approved templates
Business definitions	Inferred from prompt/schema names	Declared in semantic layer
Unsupported requests	Often answered by nearest plausible SQL	Rejected or clarified when outside coverage
Dashboard output	Usually one-off query result	Saved, refreshable widget artifact
Safety basis	Prompt instruction and filtering	No model-authored SQL execution path

5.7 Interpretation

The evaluation supports three conclusions. First, a semantic-layer architecture can cover a broad set of natural e-commerce analytics questions without exposing users to raw SQL. Second, refusal is part of correctness: a system that answers an unsupported question with a plausible but wrong query is less trustworthy than one that declines. Third, the remaining Admin fidelity gap is an implementation gap in the current compiler, not evidence that the architecture must be replaced.

This interpretation is consistent with mixed-initiative NLIDB work such as [5], which showed that users do not reliably detect wrong system answers without help. In a reporting context, a visible refusal or clarification is therefore often safer than a confident wrong chart.

CHAPTER 6

LIMITATIONS AND FUTURE WORK

6.1 Limitations

AEGIS is intentionally bounded. Its safety and auditability come from the fact that all answerable concepts must be declared in the semantic layer and all executable SQL must be produced by deterministic compiler templates. The limitations below should therefore be read as explicit boundaries of the current nopCommerce prototype, not as reasons to bypass the architecture with free-form SQL generation.

- **Single-domain evaluation:** The final evaluation is over one e-commerce deployment, nopCommerce. The results show that the architecture works in this domain, but they do not prove cross-domain generality.
- **Author-generated natural-language data:** The 500-question dataset is static and checkable, but it is still author-generated. A stronger study would collect questions from store owners or administrators and annotate them with at least two reviewers.
- **Finite semantic coverage:** Boundary questions about web telemetry, marketing attribution, support tickets, review-text sentiment, forecasting, churn prediction, supplier performance, fraud scoring, delivery SLA analysis, and product affinity are deliberately outside the current semantic layer.
- **Remaining Admin fidelity gap:** The Admin oracle benchmark reaches 15/16 result accuracy. The remaining dashboard mismatch requires a general multi-period matrix-summary primitive, not a hardcoded report-name preset.
- **LLM dependence for intent extraction:** The compiler and safety layer are deterministic, but natural-language intent extraction still depends on the configured LLM API. Future work should compare multiple OpenAI-compatible models on the same static dataset.
- **Prototype database target:** The evaluated prototype targets nopCommerce on MySQL. This is an implementation and evaluation-scope choice, not an architectural limitation; other databases require dialect-specific compiler templates and safety rules.

- **Widget persistence:** The prototype stores widget metadata in simple local persistence. A production deployment should move the widget registry to a transactional database with migrations, ownership policies, and administrative audit views.

6.2 Future Work

- **Matrix summaries:** Add the general matrix-summary primitive needed for the remaining Admin fidelity mismatch, while keeping the compiler template-based and avoiding report-specific presets.
- **Cross-schema evaluation:** Repeat the static-dataset process on a second schema such as WooCommerce or a non-commerce operational database.
- **Semantic-layer tooling:** Build a guided interface so analysts can define metrics, dimensions, predicates, join paths, and display labels without editing Python code.
- **Human dataset study:** Collect real user questions and perform independent answerability and oracle annotation, including inter-annotator agreement.
- **Model comparison:** Compare parser accuracy, refusal behavior, and execution validity across several OpenAI-compatible LLM APIs while keeping the compiler and semantic layer fixed.

CHAPTER 7

CONCLUSION

This thesis presented AEGIS, a constraint-based architecture for safe LLM-assisted natural-language analytics over relational databases. The system limits the language model to intent extraction while query construction, chart selection, and widget persistence are handled by deterministic components. This design makes approved business terms explicit through a semantic layer and avoids exposing raw SQL generation authority to the language model.

The final evaluation used static nopCommerce datasets rather than an ad hoc one-time question set. On the 500-question live benchmark, AEGIS parsed 498 of 500 prompts, answered and executed 422 of 425 supported requests, and rejected or clarified 74 of 75 realistic boundary requests. Against 16 source-derived Admin analytics oracles, it achieved 16 of 16 execution validity, 16 of 16 shape accuracy, and 15 of 16 result accuracy. The focused semantic-coverage suite further confirmed representative supported and boundary cases.

These results should be read with the right scope. AEGIS is not an infinite natural-language-to-SQL engine, and the thesis does not claim perfect accuracy. Its contribution is a bounded reporting architecture: if the semantic layer defines a business concept, the system can map natural language to a safe, refreshable analytical widget; if the concept is outside the layer, the correct behavior is to reject or clarify rather than invent an answer.

The remaining Admin mismatch strengthens the evaluation because it shows that the benchmark can expose real implementation gaps. Fixing that gap should extend the general compiler with a reusable matrix-summary primitive, not add a nopCommerce-specific shortcut. Within this scope, AEGIS demonstrates a practical path toward safer, more auditable natural-language analytics in institutional dashboard systems.

REFERENCES

- [1] T. Yu et al., "Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-SQL task," in Proc. 2018 Conf. Empirical Methods in Natural Language Processing (EMNLP), 2018, pp. 3911-3921.
- [2] J. Li et al., "Can LLM already serve as a database interface? A big bench for large-scale database grounded text-to-SQLs," in Advances in Neural Information Processing Systems (NeurIPS), vol. 36, 2023.
- [3] K. Affolter, K. Stockinger, and A. Bernstein, "A comparative survey of recent natural language interfaces for databases," The VLDB Journal, vol. 28, no. 5, pp. 793-819, 2019.
- [4] M. Liu and J. Xu, "NLI4DB: A systematic review of natural language interfaces for databases," arXiv:2503.02435, 2025.
- [5] F. Li and H. V. Jagadish, "Constructing an interactive natural language interface for relational databases," Proceedings of the VLDB Endowment, vol. 8, no. 1, pp. 73-84, 2014.
- [6] C. Lehmann, D. Gehrig, S. Holdener, C. Saladin, J. P. Monteiro, and K. Stockinger, "Building natural language interfaces for databases in practice," in Proc. 34th Int. Conf. Scientific and Statistical Database Management (SSDBM), 2022, Art. no. 20.
- [7] B. Wang, R. Shin, X. Liu, O. Polozov, and M. Richardson, "RAT-SQL: Relation-aware schema encoding and linking for text-to-SQL parsers," in Proc. 58th Annual Meeting of the Association for Computational Linguistics (ACL), 2020, pp. 7567-7578.
- [8] T. Scholak, N. Schucher, and D. Bahdanau, "PICARD: Parsing incrementally for constrained auto-regressive decoding from language models," in Proc. 2021 Conf. Empirical Methods in Natural Language Processing (EMNLP), 2021, pp. 9895-9901.
- [9] H. S. Shalaan, T. H. A. Soliman, and A. M. Abdelaziz, "G-SQL: A schema-aware and rule-guided approach for robust natural language to SQL translation," IEEE Access, vol. 13, pp. 158520-158534, 2025.
- [10] X. Su, Y. Gu, P. Wang, W. Gu, L. Qi, and J. He, "A robust natural language text-to-SQL generation framework with dynamic strategies based on LLMs," Scientific Reports, vol. 16, Art. no. 7892, 2026.
- [11] A. Narechania, A. Srinivasan, and J. Stasko, "nl4dv: A toolkit for generating analytic specifications for data visualization from natural language queries," IEEE Transactions on Visualization and Computer Graphics, vol. 27, no. 2, pp. 369-379, 2021.

- [12] T. Gao, M. Dontcheva, E. Adar, Z. Liu, and K. G. Karahalios, "DataTone: Managing ambiguity in natural language interfaces for data visualization," in Proc. 28th Annual ACM Symposium on User Interface Software and Technology (UIST), 2015, pp. 489-500.
- [13] D. Deng, A. Wu, H. Qu, and Y. Wu, "DashBot: Insight-driven dashboard generation based on deep reinforcement learning," IEEE Transactions on Visualization and Computer Graphics, vol. 29, no. 1, pp. 690-700, 2023.
- [14] G. N. Shailesh, M. Pavithran, A. Rahul Hari Venkat, and P. Kaliappan, "Conversational BI: Natural language interface to business dashboards," International Journal of Engineering Research & Technology, vol. 14, no. 12, 2025.
- [15] J. Valkenburgh, "Enhancing business dashboards with explanatory analytics and AI: Exploring the use of AI and explanatory analytics to enhance business decision-making," M.S. thesis, Information Management, Turku School of Economics, University of Turku, Finland, 2024.
- [16] OpenAI, "Introducing structured outputs in the API," OpenAI, 2024.